ChatGPT

OpenAI’s Blueprint to Power the Future Web with AI

Introduction

OpenAI has rapidly evolved from a little-known research lab into a central player of the AI boom [1]. In just three years since launching ChatGPT, the company is no longer content to be seen as a simple chatbot maker. It’s charting a bold plan to become a **foundational infrastructure layer** for the next generation of web applications, using large language models like ChatGPT to transform how we interact online [2] [3]. Recent moves – from integrating third-party apps into ChatGPT to unveiling new developer tools – all signal OpenAI’s ambition to remake the web into an AI-driven, conversational experience. The following analysis examines OpenAI’s strategy and the tools it's rolling out, and what this could mean for users, developers, and businesses in the near future.

ChatGPT: From Chatbot to AI Platform

ChatGPT’s mobile app has become a gateway to many services, reflecting OpenAI’s vision of a universal conversational interface. The company announced that users will soon be able to get things done online “without ever having to leave ChatGPT,” performing tasks like creating Spotify playlists or browsing Zillow listings right within a chat [3]. By incorporating such third-party apps directly into its conversations, ChatGPT is positioning itself much like an operating system for the web – analogous to how Apple’s iOS or Google’s Android serve as central hubs for mobile apps [3].

OpenAI’s CEO Sam Altman has described ChatGPT’s expansion as moving beyond a single app into a **platform** spanning shopping, entertainment, education, and even government services [2]. The goal is to put ChatGPT at “the core of how people use technology” in the AI era [3]. In practice, this means ChatGPT is gaining capabilities that mirror or integrate with many traditional websites and apps. For example, ChatGPT can now plan trips, write small programs, browse the web on a user’s behalf, and even execute purchases through a new “instant checkout” feature [3] [4]. Instead of hopping between different websites or mobile apps, a user could simply **ask ChatGPT** to handle tasks end-to-end – whether it’s ordering a meal, playing a song, or finding an apartment listing – and see the results or interactive widgets right in the chat window. This represents a significant shift in the digital interface: from clicking and tapping in separate apps, to conversing naturally with one AI assistant that orchestrates many services behind the scenes.

Critically, ChatGPT’s massive user base gives OpenAI a strong springboard for this platform vision. The chatbot now boasts **800 million weekly active users**, an astonishing scale reached in under three years [5]. That kind of reach rivals the biggest social networks and far exceeds any single web service’s user count. OpenAI is clearly leveraging this popularity: by inviting external apps into ChatGPT, it hopes to lock in users who find it convenient to accomplish everything in one place. Tech analysts have begun likening ChatGPT to an emerging “AI operating system” because of these integrated capabilities [6]. In essence, OpenAI wants ChatGPT to be the new interface for the web – a **conversational layer** that sits on top of traditional websites and apps, handling user requests with AI and automating the underlying web tasks. If this vision succeeds, it could redefine how we access information and services online.

Building the AI Web: OpenAI as Infrastructure

OpenAI's grand ambitions go beyond software. To truly become the backbone of an AI-powered web, OpenAI is investing heavily in **physical infrastructure** – namely the data centers and hardware needed to run powerful AI models at scale. Sam Altman's company is reportedly "shelling out billions" on this front [7]. In fact, at its recent developer conference, OpenAI announced plans to spend enormous sums on new data center capacity, including a deal for **6 gigawatts** of cloud computing power using AMD chips [7]. This comes on top of multi-billion-dollar agreements with Nvidia (to secure cutting-edge AI GPUs) and with Oracle, its cloud partner, to expand global hosting for OpenAI's models [7] [8].

OpenAI's Stargate project exemplifies this infrastructure push. The company partnered with Oracle and others to invest up to $500 billion in new AI data centers, such as this large facility under construction in Texas [9]. According to OpenAI President Greg Brockman, "we need as much computing power as we can possibly get" to support advanced AI services [9]. By building out massive server farms around the world, OpenAI is trying to ensure it has the raw horsepower to run complex AI applications for hundreds of millions of users simultaneously.

This strategy is reminiscent of the playbooks of earlier tech giants. Just as Google built an empire on a vast network of data centers to index and serve the web, OpenAI is now spending aggressively to become a **core infrastructure provider** for the AI era [10] [11]. The company's investments include a partnership (with Oracle and SoftBank) to construct multiple AI supercomputing hubs in the U.S. under the **Stargate** initiative [12], as well as agreements to develop AI infrastructure in the U.K. and Middle East [13]. In total, OpenAI's commitments to hardware are staggering – from a $100 billion deal with Nvidia for chips to a $300 billion cloud spend with Oracle [14].

Why such a massive outlay? OpenAI faces a **chicken-and-egg** challenge [15]. Its platform needs powerful computing to deliver advanced AI features, but scaling up computing is extremely expensive. Meanwhile, to fund those expenses, OpenAI needs revenue from users and developers adopting its platform [16]. So the company is attempting to grow both sides in tandem: launch new features and an app ecosystem to attract users (and paying customers), while pouring money into infrastructure to handle the growth [16]. It's a high-stakes bet – OpenAI is *not yet profitable* and reportedly lost $7.8 billion in the first half of 2025 [17]. Nonetheless, this aggressive expansion is intended to make OpenAI's technology **indispensable**. As one analyst put it, today it's "ChatGPT versus the Microsoft bundle" – meaning OpenAI must compete with tech giants that can embed AI across their entrenched products [18]. By becoming the foundational layer for many applications (rather than just one chatbot), OpenAI hopes to embed itself deeply into the fabric of the web, much like Amazon did with cloud computing or Microsoft with the PC software ecosystem.

New Tools Fueling an AI Ecosystem

To achieve its platform vision, OpenAI is actively courting developers with a suite of new tools and SDKs announced in late 2025. These tools are meant to make it easier to build AI-driven apps and **autonomous agents** on top of OpenAI's models, thereby creating a robust developer ecosystem. At OpenAI's DevDay 2025 conference, Altman unveiled several key components that expand what developers (and by extension, users) can do with ChatGPT [19] [20]. The major announcements included:

- **Apps SDK:** A software kit that lets developers create **third-party applications to run natively inside ChatGPT**. Rather than redirecting users to external websites or apps, an Apps SDK app can present interactive content (cards, forms, images, etc.) directly within the chat interface [21]. This is a game-changer because it transforms ChatGPT into a platform akin to an app store or

web browser – users can invoke services like Spotify, Expedia, or Zillow conversationally and get rich results inline [22]. Initial partners like Booking.com, Canva, and Spotify built plug-in apps so that, for example, a user can say "Spotify, make a playlist for my party" and get a playable playlist without leaving ChatGPT [22]. The Apps SDK provides the protocols (such as the new Model Context Protocol) for these apps to exchange data with ChatGPT securely and contextually. In short, it **embeds the web's functionality into ChatGPT**, fulfilling OpenAI's goal of an all-in-one experience.

- **AgentKit:** A complete toolkit for developers to **build, deploy, and optimize AI agents**. OpenAI describes AgentKit as a "set of building blocks" to take an AI agent from prototype to production with minimal friction [19]. In practice, AgentKit streamlines the previously complicated process of creating agents that can take actions, use tools, and perform multi-step tasks autonomously. Its components include a visual **Agent Builder** (a drag-and-drop interface to design agent logic and workflows), **ChatKit** (an embeddable chat UI component for plugging an agent into any app or website), and a **Connector Registry** for integrating external tools and data sources securely [23] [24]. By providing these pieces out-of-the-box, AgentKit lowers the barrier for developers to create sophisticated agents – whether it's a customer service bot, an AI research assistant, or an automated task manager. Sam Altman likened the Agent Builder to "Canva for building agents," underscoring its ease of use [25]. Overall, AgentKit is about **making AI agents mainstream** by giving developers a one-stop solution to build reliable agents that can act on a user's behalf, not just chat passively [26].

- **Agents SDK:** A programming framework introduced earlier in 2025 to support agent development. The Agents SDK is essentially the **underlying orchestration layer** that AgentKit builds upon [27]. It provides developers with the APIs and libraries to manage complex agent behaviors, such as calling tools, handling multi-agent collaborations, maintaining session memory, and enforcing safety constraints. In OpenAI's words, the Agents SDK allows developers to "orchestrate single-agent and multi-agent workflows," abstracting away much of the heavy lifting involved in coordinating how an agent thinks, decides, and interacts with external systems [28]. For example, it can handle an "agent loop" where the AI model iteratively reasons, uses a tool (like a web search), evaluates the result, and continues until a goal is achieved [29] [30]. By packaging best practices (like guardrails for safety and traceability for debugging) into a ready-made framework, the Agents SDK ensures that anyone building on OpenAI's platform can create **autonomous agents with far less effort and greater reliability** than before. It complements AgentKit's higher-level tools by giving more technical developers fine-grained control when needed.

Together, these components lay the groundwork for a **new kind of developer ecosystem** centered on AI agents and ChatGPT-based applications. In effect, OpenAI is providing the analogous pieces of an app platform: a place to run apps (ChatGPT with Apps SDK), tools to build complex app logic (AgentKit and Agents SDK), and even distribution and monetization channels (an upcoming app directory and usage-based revenue models). This mirrors what successful platforms like Windows, iOS, or cloud providers have done – but tuned to the AI age. By making ChatGPT extensible and **developer-friendly**, OpenAI is inviting programmers and companies to build on its infrastructure rather than creating competing AI systems. As Altman emphasized, these tools are meant to reduce friction and "help you take agents from prototype to production… with way less friction" [19] [31]. The bet is that an army of third-party developers will extend ChatGPT's capabilities in ways OpenAI alone might not imagine, driving a virtuous cycle: more useful apps and agents attract more users, which in turn attracts more developers to the platform.

Implications for Users: A New Interface for the Web

For everyday internet users, OpenAI's platform push could bring about a **fundamental change in digital interactions**. The vision is that instead of using dozens of separate apps and websites through graphical interfaces, users will increasingly accomplish tasks through a **single conversational AI interface**. This doesn't mean apps and sites disappear; rather, ChatGPT will act as an intelligent broker, handling the dialogue with users and then invoking the appropriate service behind the scenes. From a user's perspective, it's like having a super-smart personal assistant who can navigate the web's offerings on command.

One immediate implication is **convenience and efficiency**. With apps like Spotify, Expedia, or Canva integrated into ChatGPT, a user can get things done with simple requests in natural language. For example, planning a trip might involve asking, "Find me a cheap flight to Paris and a hotel near the Eiffel Tower next month," and the AI will consult a travel app, compare options, and present the best choices in-chat. This eliminates the tedious process of manually visiting multiple websites or apps and piecing together information. OpenAI explicitly wants users to "get things done online without ever having to leave ChatGPT" [3], which suggests a seamless experience where the boundaries between different services are invisible. Digital interfaces could become **more conversational and personalized**, with ChatGPT remembering context from your previous interactions (your preferences, past orders, schedule, etc.) to tailor its assistance. This is a step beyond today's voice assistants or chatbots, which are often siloed to one service; OpenAI's approach aims to unify many services under one AI assistant that truly works across domains.

Another implication is how users discover and use new services. In a ChatGPT-centric web, you might not need to install new apps or sign up on new websites yourself – the AI will bring the service to you as needed. For instance, if you express a desire to learn a new skill, ChatGPT might automatically surface a Coursera app or a recommended educational tool within the conversation. This **model-driven discovery** means the AI can suggest apps contextually (with your permission) when they are relevant to your request [32] [22]. It's akin to having a smart web navigator that knows which tool is right for the job and calls it instantly. For users, this could reduce the clutter of having many apps installed and accounts everywhere. The trade-off is putting a lot of trust in ChatGPT to mediate your online activities. Questions of reliability, privacy, and bias will naturally arise – users will need confidence that the AI is making good choices and not favoring certain providers unfairly. Nonetheless, if OpenAI and its partners execute well, general users stand to benefit from **AI-assisted simplicity**: no more juggling passwords, app stores, or search queries for routine tasks, just a conversational partner that handles it all.

Finally, the user experience online could become more **proactive and interactive** thanks to AI agents. Instead of just answering questions, ChatGPT (with agents built via AgentKit) could take initiative on your behalf. Imagine telling an agent to "monitor my deliveries and reschedule any that conflict with my calendar" – the agent could continuously work in the background without further prompting. Or a study-planner agent might proactively suggest a learning schedule each week based on your progress. In essence, OpenAI's platform hints at everyday users having access to **autonomous digital assistants** that don't just respond, but can act and manage tasks. This blurs the line between using a service and delegating a task to an AI. The web could feel more tailored: interfaces that adapt to your needs in real-time, possibly even generating custom content (like short-form videos via AI, as OpenAI's Sora app does [33]). For the average person, this means the web becomes less about navigating menus or forms, and more about simply stating goals and receiving outcomes. It's a transformative idea of human-computer interaction – if widely adopted, chatting with an AI agent could become as common as tapping on apps is today.

Implications for Developers: A New AI Marketplace

Developers are central to OpenAI's strategy, and the company's recent moves create a host of **new opportunities** – as well as new considerations – for those building applications. OpenAI is effectively launching a **platform ecosystem** that invites developers to create apps and agents that plug into ChatGPT, similar to how Apple's App Store or Google's Play Store enabled third parties to expand the capabilities of those platforms. The introduction of the Apps SDK and AgentKit means that a developer can build an AI-powered service and instantly reach ChatGPT's vast user base without needing to build a standalone website or mobile app from scratch. Sam Altman highlighted this distribution potential, noting that an app built for ChatGPT could tap into "hundreds of millions" of users on day one [34]. For developers and startups, this is a compelling proposition: **write one ChatGPT app and it's available to a massive audience** that is actively engaging with AI for all sorts of tasks. In many ways, OpenAI is offering to handle the "last mile" of user interaction (via the chat interface) and even discovery of your app, letting developers focus on their service's core logic.

The platform comes with a suite of **powerful capabilities** that developers can leverage. Thanks to the Agents SDK and built-in tools, an independent developer can now easily give their app abilities like web browsing, file search, or even executing code, all via OpenAI's APIs [35] [28]. This drastically lowers the complexity versus building or integrating those functions on their own. Moreover, AgentKit's visual workflow designer and connectors mean that a small team could create fairly complex agent behaviors (for example, an agent that checks inventory in a database, then drafts an email to order new stock when it's low) without needing deep machine learning expertise. Essentially, OpenAI's platform **abstracts a lot of AI complexity** into accessible tools. Developers can mix and match OpenAI's language models with their own business logic and data, yielding custom agents or apps that would have been very hard to build from scratch. Early adopters have reported dramatic productivity gains – for instance, the software company Canva built a customer support agent in under two hours using the AgentKit visual builder [36]. Such speed and ease suggest that the barrier to entry for AI development is coming down, enabling more developers (even those not specialized in AI) to participate in building next-gen applications.

However, developers must also consider the **platform dynamics**. By building on ChatGPT, they are to some extent **entrusting their product's distribution and user experience to OpenAI's ecosystem**. This is analogous to developing for Facebook or the iOS App Store; the platform can provide tremendous reach, but it also sets the rules (review policies, revenue sharing, technical constraints). OpenAI has indicated an app review system and directory will launch for ChatGPT apps [37], likely meaning developers will have to meet certain guidelines for their apps to be approved and discoverable. There's also competition: as more developers flock to create ChatGPT apps, standing out could be a challenge, and the assistant might favor the best-designed or most relevant apps in its suggestions. On the flip side, OpenAI's integrated approach might reduce the "winner-takes-all" effect seen in traditional app stores, since ChatGPT could flexibly recommend niche apps whenever appropriate to a user's query. Importantly, OpenAI's move may spur **competition with other AI platforms**. Companies like Google, Microsoft, and startups like Anthropic are also courting developers to build on their AI models and assistants. OpenAI's advantage is its head start with ChatGPT's popularity and a relatively neutral platform (since OpenAI isn't also pushing a huge existing ecosystem of its own non-AI products, aside from ChatGPT itself). If it can keep developers happy with favorable economics and powerful tools, OpenAI could become to AI applications what Apple became to mobile apps or what Amazon Web Services became to cloud apps – the default platform. As one analysis noted, OpenAI is "rolling out the red carpet" for developers with unprecedented tools and channels to monetize AI experiences [38]. In summary, for developers this moment looks like the **dawn of a new marketplace**: those who get in early and build high-quality ChatGPT apps or agents have a chance to

ride the wave of AI-native "killer apps" in the coming years, much like the early days of smartphone apps or cloud services – albeit under the governance of OpenAI's ecosystem.

Implications for Businesses: Automation and Integration

Business leaders are eyeing OpenAI's evolving platform both as a competitive threat and as a promising new tool for innovation. On one hand, if ChatGPT and its agent ecosystem become a primary way consumers access services, companies will need to **adapt their strategies to remain visible and relevant**. Just as businesses in the 2010s scrambled to build mobile-friendly sites or native apps, today they may need to ensure they have a presence in conversational AI platforms. This could mean developing a ChatGPT app that interfaces with their product (e.g. a retail chain making a shopping assistant agent, or a bank creating a finance coach agent) so that when users ask the AI for something in that domain, their service is in the mix. Being an early partner, as brands like Spotify and Expedia were at launch, can yield a first-mover advantage in shaping consumer expectations in this new channel [22]. Companies will also have to consider **data integration** – OpenAI's Connector Registry and tools allow agents to tie into internal company systems securely [23]. For enterprise leaders, this presents an opportunity to streamline operations: for example, connecting an AI agent to inventory databases, CRM systems, or scheduling tools to handle routine queries or tasks that currently require human coordination.

Another huge implication is the potential for **automation of complex workflows**. OpenAI's agents are not just chatbots answering FAQs; they can perform multi-step operations, call external APIs, and make decisions towards a goal. For businesses, this opens the door to automating processes that involve unstructured interactions or multiple systems – areas that historically have been hard to fully automate. Consider customer support: with AgentKit, a company could deploy an agent that not only answers a customer's question, but also pulls up their account info, makes a change the customer requested, and drafts a follow-up email, all in one conversational flow. Early examples show this is feasible: HubSpot, for instance, was able to quickly set up a customer support agent using these tools [36]. Over time, such AI agents could handle a large chunk of routine inquiries or tasks, freeing human staff to focus on more complex or high-touch issues. In fields like finance, law, or medicine, we might see AI agents assisting professionals by doing first-pass analysis or paperwork (under human supervision). For **business leaders**, the takeaway is that AI is moving from a novelty to a practical *workflow tool*. Integrating OpenAI's capabilities could reduce costs (through automation of labor-intensive tasks) and enable new services (like 24/7 personalized customer interactions) that were not economically viable before.

Of course, this revolution comes with challenges. Companies will need to address **trust and oversight** of AI actions: an agent acting on behalf of your business must behave reliably and in line with policies. OpenAI's platform provides guardrails and evaluation tools to help with this [39] [40], but it remains the business's responsibility to supervise and fine-tune their AI agents. There are also strategic considerations: relying on OpenAI's ecosystem means ceding some control (and potentially data) to an external platform. Some enterprises might hesitate to let ChatGPT sit between them and their customers. We may see larger players pursue a dual strategy – integrating with ChatGPT for broad reach, while also developing their own AI systems for more control (like banks doing both open platform chatbots and proprietary AI co-pilots internally). Nonetheless, ignoring the trend is risky. As one venture investor noted, OpenAI is trying a "spend aggressively to become essential, monetize later" approach that has worked for tech giants before [41]. If it succeeds, businesses that failed to engage with these AI platforms might find themselves **disintermediated** – imagine a future where consumers simply ask an AI to "book me the best insurance plan" and the AI's chosen app (maybe not the incumbent insurer's) handles everything. The businesses that thrive will likely be those that **embrace AI integration** early: using tools like the Agents SDK to enhance their operations, and ensuring their

services can be plugged into conversational interfaces where the next generation of customers will be active.

Looking Ahead: An AI-First Web

OpenAI's recent maneuvers paint a picture of a future web that is **AI-first and user-centric** in new ways. The company's plan – to serve as the backbone for AI applications and experiences across the internet – is immensely ambitious. It involves technical innovation (scaling up powerful models and infrastructure), ecosystem building (attracting developers and partners), and shaping user behavior (getting people comfortable with an AI as the go-to intermediary for tasks). In many respects, OpenAI is attempting to do in a few years what platforms like Windows, Android, or AWS did over decades: become the default layer that others build on. The difference now is the **speed and scale** of AI adoption. ChatGPT's explosive growth to 800 million weekly users hints that if OpenAI provides compelling functionality, people will flock to it [5].

The next couple of years will be critical in determining whether OpenAI truly becomes the "foundation of the future web" or whether this vision splinters. We can expect competition to intensify – rival tech firms are not standing still. Google is weaving AI into its search and productivity tools; Meta is open-sourcing AI models; Microsoft is deeply integrating GPT-4 into its Windows and Office platforms. OpenAI, lacking a traditional hardware or OS ecosystem, is effectively trying to **leapfrog** into that role with ChatGPT as a new kind of universal platform. This could lead to partnerships (or conflicts) with the incumbents. Notably, OpenAI has hinted at hardware projects (like a mysterious device with former Apple designer Jony Ive) that could further entwine its AI with consumer life [42].

For general users, if OpenAI's bet pays off, the web in five years might feel drastically different. We might converse our way through most online interactions, with AI agents handling everything from daily errands to creative projects. Websites and apps might evolve to be *AI-accessible* modules rather than full-fledged destinations, as traffic shifts to AI intermediaries. For developers and businesses, an OpenAI-dominated ecosystem would require aligning with its platform to reach consumers, much as aligning with mobile app stores became necessary in the smartphone era. There are risks – centralization of so much activity on one AI platform raises concerns around bias, moderation, and single points of failure (indeed, ChatGPT outages have already caused widespread disruption). OpenAI will have to address these as it scales.

In summary, OpenAI's plan to power the future of the web is a **bold, high-risk, high-reward strategy**. The company is pouring resources into becoming both the **brain** (through advanced AI models) and the **backbone** (through infrastructure and developer frameworks) of a new web paradigm. It's a play to redefine how we use technology – making interaction more natural and unified via AI – while capturing the value that comes from being the central platform that enables it all. As history has shown with previous tech shifts, the winners are often those who create the most useful platforms and nurture vibrant ecosystems around them. OpenAI is clearly betting that an AI-centric platform will be the next such winner. If it succeeds, we may look back on this period as the moment the web entered a new phase – one where **chatbots became the new browsers, AI agents the new apps, and cloud-scale AI the new infrastructure of everyday life.**

Sources: OpenAI/CNN Business [3] [43] [9] [8]; TechCrunch [19] [20]; TokenRing AI [22] [23]; Indian Express [44] [45]; OpenAI Blog [28]; Data Science Dojo [46].

---

1 2 3 4 5 7 8 9 10 11 12 13 14 15 16 17 18 33 41 42 43 OpenAI wants to build the next era of the web, and it’s shelling out billions to do it - Egypt Independent
https://www.egyptindependent.com/openai-wants-to-build-the-next-era-of-the-web-and-its-shelling-out-billions-to-do-it/

6 21 22 23 27 38 OpenAI DevDay 2025: ChatGPT Transforms into an AI Operating System with Apps SDK, AgentKit, and More
https://markets.financialcontent.com/wral/article/tokenring-2025-10-6-openai-devday-2025-chatgpt-transforms-into-an-ai-operating-system-with-apps-sdk-agentkit-and-more

19 20 24 25 26 31 OpenAI launches AgentKit to help developers build and ship AI agents | TechCrunch
https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/

28 35 40 New tools for building agents | OpenAI
https://openai.com/index/new-tools-for-building-agents/

29 30 46 OpenAI Agents SDK: The Modern Guide to Agentic AI Development | Data Science Dojo
https://datasciencedojo.com/blog/openai-agents-sdk/

32 39 44 45 OpenAI DevDay 2025: ChatGPT gets apps, AgentKit for developers, and cheaper GPT models
https://indianexpress.com/article/technology/artificial-intelligence/openai-devday-2025-chatgpt-gets-apps-agentkit-for-developers-and-cheaper-gpt-models-10292443/

34 36 37 Everything OpenAI Released on DevDay 2025, Explained | The Neuron
https://www.theneuron.ai/explainer-articles/everything-openai-released-on-devday-2025-explained